

Constraint-based Scheduling & Packing including MiniBrass



Agenda



Grundlagen:

- Constraint Satisfaction (Optimisation) Problems
- Funktionsweise von Constraint-Lösern und Sprachen (MiniZinc)
- Einsatz im Scheduling

ISSE-Entwicklungen:

- Constraint Relationships für Soft Constraints
- Entwickelte Fallstudien
- Sprachunterstützung / Features

Constraint Programming: Einordnung



- Generischer Ansatz zur Lösung von Erfüllbarkeitsproblemen
- Ausnützen von Struktur von logischen Bedingungen (Tsang, 1993)
- Konzentration auf endliche Wertebereiche und Erfüllbarkeit (Russell and Norvig, 2010, Kap. 5)
 - Im Gegensatz z.B. zu linearer Programmierung, konvexe Optimierung
 - Verallgemeinert Boolesche Erfüllbarkeitsprobleme (SAT)
 - Scheduling-Probleme
 - Reorganisationen
 - Zuweisungsprobleme (z.B. Frequenzen an Sender, Energie and Produzenten, ...)
- Deklarativ, aber Constraints sind an Algorithmen geknüpft (Rossi et al., 2006)!

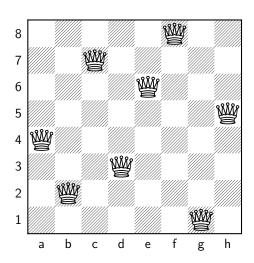
Ein prototypisches CSP





Und noch eins







Definition (Constraint-Problem)

Ein Constraint-Problem (X, D, C) ist beschrieben durch

• Variablen X, Domänen $D = (D_x)_{x \in X}$ (endlich), Constraints C

Domänen typischerweise: int, bool, (float)

Was ist ein **Constraint**? Ein *boolesches* Prädikat über einer Belegung von X: $X = \{x, y, z\}$

- x < y
- x + 5 = z y
- alldifferent([x, y, z])
- ∀ und ∃ nur über endlichen Wertebereichen



Problem

CSP(X, D, C) mit

- $X = \{x, y, z\}$
- $D_x = D_y = \{0, 1, 2\}, D_z = \{0, 1\}$
- 0
- $c_1: x \neq y, y \neq z, x \neq z$
- $c_2: x+1=y$

```
var 0..2: x;
var 0..2: y;
var 0..1: z;

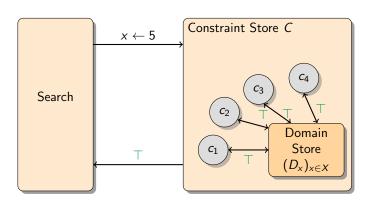
% c1
constraint x != y /\ y != z /\ x !=
% c2
constraint x + 1 = y;
solve satisfy;
```

Welche Zuweisung ist eine Lösung dieses Problems?

- $\Theta = \{(x \to 1, y \to 2, z \to ?), (x \to 0, y \to 1, z \to ?)\}$ erfüllen c_2 ;
- $(x \to 0, y \to 1, z \to ?)$ lässt sich aber zu keiner Lösung erweitern, da z entweder 0 oder 1 sein muss und somit garantiert c_1 verletzt
- Also ist die einzige Lösung $(x \to 1, y \to 2, z \to 0)$

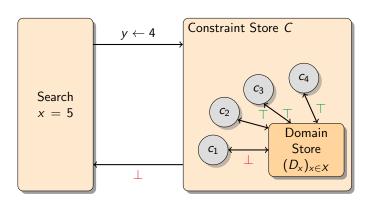
Architektur von Constraint-Lösern





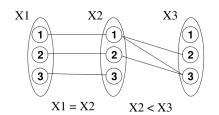
Architektur von Constraint-Lösern

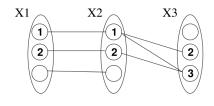




Constraint-Propagation: Beispiel







Entfernen von Werten, die zu keiner Lösung führen können.

Globale Constraints



- Betrachten wir folgendes einfaches Problem
 - $X = \{x_1, x_2, x_3\}$
 - $(D_x)_{x \in X} = \{1, 2\}$
 - $C: x_1 \neq x_2, x_2 \neq x_3, x_1 \neq x_3$
- Ist dieses Problem nach Constraint-Propagation mit binären Constraints lösbar?
- Ja, für jedes $d \in D_x$ gibt es einen Partner
- Insgesamt allerdings nicht, da mindestens 3 unterschiedliche Werte nötig
- $\bullet \to \mathsf{daher}$ globale Constraints, die eine größere Menge von Variablen im Auge betrachten können
- Und spezialisierte Propagationsalgorithmen haben!
- all different (x_1, x_2, x_3)

Task-Zuweisung in Practice



- Taskzuweisungsproblem (task allocation problem)
 - n Roboter
 - m Tasks
 - Gebe jedem Roboter einen unterschiedlichen Task, um den Gewinn zu maximieren (Unterschied zu Kompensation und Strafe)
- Beispielproblem:

$$- n = 4, m = 5$$

	t1	t2	t3	t4	t5
r1	7	1	3	4	6
r2	8	2	5	1	4
r3	4	3	7	2	5
r4	3	1	6	3	6

Task-Zuweisung: Modell



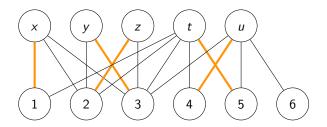
```
% problem data
int: n; set of int: ROBOTS = 1..n;
int: m; set of int: TASKS = 1..m;
array[ROBOTS,TASKS] of int: profit;
% decisions
array[ROBOTS] of var TASKS: allocation;
% goal
solve maximize sum(r in ROBOTS) (profit[r, allocation[r]] );
% have robots work on different tasks
constraint alldifferent(allocation);
```

AllDifferent – Machbarkeit



- Erster bekannter Propagator
- Basiert auf Matching in bipartiten Graphen
- Laufzeit ist polynomiell!

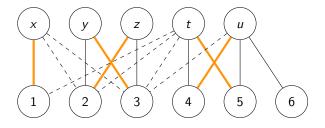
```
var {1,2,3}: x; var {2,3}: y; var {2,3}: z;
var {1,2,3,4,5}: t; var {3,4,5,6}: u;
constraint alldifferent([x,y,z,t,u]);
```



AllDifferent - Propagierung



```
var {1,2,3}: x; var {2,3}: y; var {2,3}: z;
var {1,2,3,4,5}: t; var {3,4,5,6}: u;
constraint alldifferent([x,y,z,t,u]);
```

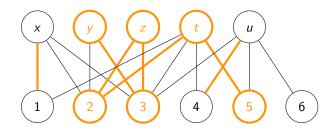


```
var {1}: x;     var {2,3}: y;     var {2,3}: z;
var {4,5}: t;     var {4,5,6}: u;
constraint alldifferent([x,y,z,t,u]);
```

AllDifferent - Algorithmus



```
var {1,2,3}: x; var {2,3}: y; var {2,3}: z;
var {1,2,3,4,5}: t; var {3,4,5,6}: u;
constraint alldifferent([x,y,z,t,u]);
```



Augmentierender Pfad:

$$y \rightarrow 3 \rightarrow z \rightarrow 2 \rightarrow t \rightarrow 5$$

Let that sink in



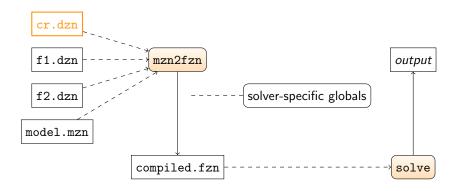
Take away message

Global constraints are the essence of constraint programming!

How can we actually use all that?



Architecture



Why MiniZinc?



Rationale

One modeling language – many solvers

Supported Solvers

- Gecode (CP)
- JaCoP (CP)
- Google Optimization Tools (CP)
- Choco (CP)
- G12 (CP/LP/MIP)



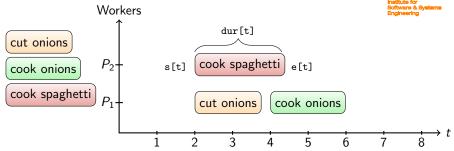
Scheduling & Packing



- A very important application for CP
- Strong support (global constraints) in MiniZinc
 - Involved propagators fortunately encapsulated by solvers!
- Planning and scheduling are coming together and constraint satisfaction may serve as a bridge between them
- Many applications in practice (production style problems)
- Overview papers:
 - Bartak'03 (Barták, 2003)
 - Bartak'10 (Barták et al., 2010)
 - Baptiste'06 (Baptiste et al., 2006)

Intro to Scheduling



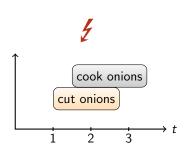


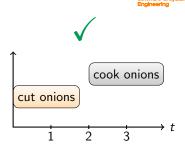
```
int: T; set of int: TASK = 1..T; set of int: HORIZON = 0..10;
array[TASK] of var HORIZON: s; % start times
array[TASK] of var HORIZON: e; % end times
array[TASK] of HORIZON: du; % durations (fixed as constants!)

constraint forall(t in TASK) (s[t] + du[t] = e[t]); % consistency
constraint e[cut_on] <= s[cook_on]; % precedence</pre>
```

Non-overlaps (Manually)

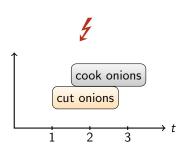


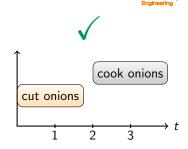




Non-overlaps (Global Constraint)

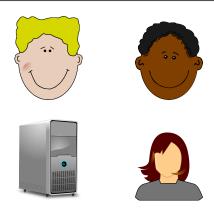






A small example





- Two brothers, Alfred and Ben; both need to
 - Read about an algorithm before
 - programming it on a computer
- Their mother needs to do work on a computer as well
- Only one computer is available



```
int: T = 5; set of int: TASK = 1...T;
set of int: HORIZON = 0..50;
array[TASK] of var HORIZON: s; % start times
array[TASK] of var HORIZON: e; % end times
var HORIZON: makeSpan;
int: readAlfred = 1; int: codeAlfred = 2;
int: readBen = 3; int: codeBen = 4; int: workMother = 5;
% durations
array[TASK] of int: d = [5, 2, 3, 8, 12];
constraint forall(t in TASK) (s[t] + d[t] = e[t]);
constraint forall(t in TASK) (e[t] <= makeSpan);</pre>
constraint e[readAlfred] <= s[codeAlfred] /\ e[readBen] <= s[codeBen]:</pre>
set of TASK: conflicting = {codeAlfred, codeBen, workMother};
constraint disjunctive([s[t] | t in conflicting],
                       [d[t] | t in conflicting] );
solve minimize makeSpan;
```

Solution

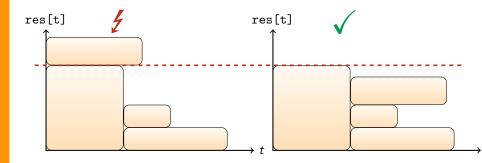


```
int: readAlfred = 1; int: codeAlfred = 2;
int: readBen = 3; int: codeBen = 4; int: workMother = 5;
set of TASK: conflicting = {codeAlfred, codeBen, workMother};
array[TASK] of int: d = [5, 2, 3, 8, 12];
s = array1d(1...5, [0, 20, 0, 12, 0]);
e = array1d(1..5, [5, 22, 3, 20, 12]);
makeSpan = 22;
         work (mother)
  read (B)
                                       code (B)
     read (A)
                                                            code (A)
```

Cumulative Scheduling



- disjunctive is useful for unary resources
- It is actually a special case of cumulative
 - I have m copies of the same resource
 - No more than *m* can be used at once!
 - Example: m workers, m bulldozers, m tools, ...
- Assume task t takes res[t] resources



Cumulative Global Constraint



MiniZinc Docs

Requires that a set of tasks given by start times s, durations d, and resource requirements r, never require more than a global resource bound b at any one time.

Household Revisited



How about buying a new computer?

Solution (1 Computer)



```
int: readAlfred = 1; int: codeAlfred = 2;
int: readBen = 3; int: codeBen = 4; int: workMother = 5;
set of TASK: conflicting = {codeAlfred, codeBen, workMother};
array[TASK] of int: d = [5, 2, 3, 8, 12];
s = array1d(1...5, [0, 20, 0, 12, 0]);
e = array1d(1..5, [5, 22, 3, 20, 12]);
makeSpan = 22;
         work (mother)
  read (B)
                                       code (B)
     read (A)
                                                            code (A)
```

Solution (2 Computers)



```
int: readAlfred = 1; int: codeAlfred = 2;
int: readBen = 3; int: codeBen = 4; int: workMother = 5;
set of TASK: conflicting = {codeAlfred, codeBen, workMother};
array[TASK] of int: d = [5, 2, 3, 8, 12];
```

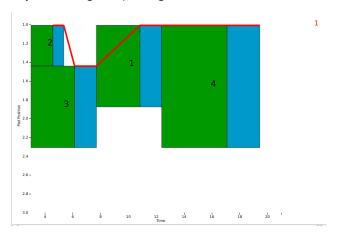
```
s = array1d(1..5 ,[0, 11, 0, 3, 0]);
e = array1d(1..5 ,[5, 13, 3, 11, 12]);
makeSpan = 13;
```

```
read (B) code (B)
read (A) code (A)
```

Port Scheduling



Combines many scheduling and packing constraints:



Scheduling: Outlook



There's clearly much more to scheduling (that could be useful for us)

- ullet Optional tasks o brings us to planning problems
- State-based models: Captures transitions of finite state machines (FSM)
- Sequence-dependent Scheduling
 - Times for cooling down
 - Color changes

Take away

Rather easy to capture since we're dealing only with constraint models!

Soft-Constraints



Präferenzen im Constraint Solving



Constraint-Problem (X, D, C)

• Variablen X, Domänen $D = (D_x)_{x \in X}$, Constraints C

In der Praxis: überbestimmte Probleme

$$\begin{aligned} & \big(\big(\{x,y,z\}, D_x = D_y = D_z = \{1,2,3\} \big), \{c_1,c_2,c_3\} \big) \text{ mit } \\ & c_1: x+1 = y \\ & c_2: z = y+2 \\ & c_3: x+y \leq 3 \end{aligned}$$

- Nicht alle Constraints können gleichzeitig erfüllt werden
 - ullet e.g., c_2 erzwingt $\mathrm{z}=3$ und $\mathrm{y}=1$, im Konflikt mit c_1
- \bullet Ein Agent wählt zwischen Zuweisungen, die $\{c_1,c_3\}$ oder $\{c_2,c_3\}$ erfüllen.

Welche Zuweisungen $v \in [X \to D]$ sollen bevorzugt werden von einem Agenten (oder sogar einer Menge von Agenten)?

Constraint Relationships



Ansatz (Schiendorfer et al., 2013)

- Definiere Relation *R* über Constraints *C* um anzugeben, welche Constraints wichtiger sind als andere, e. g.
 - c₁ wichtiger als c₂
 - ullet c_1 wichtiger als c_3



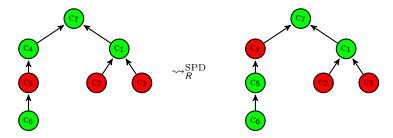
Benefits

- Qualitativer Formalismus einfach zu spezifizieren
 - Hebe diese Relation auf Verletzungsmengen
 - Dominanzeigenschaften regulieren den Tradeoff "Hierarchie vs. Egalitär"
 - Single-Predecessors-Dominance (SPD) vs. Transitive-Predecessors-Dominance (TPD)

Single-Predecessor-Dominance (SPD) Lifting

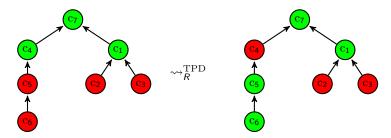


isWorseThan-Relation für Mengen verletzter Constraints



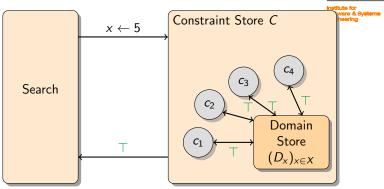
Transitive-Predecessors-Dominance (TPD) Lifting

isWorseThan-Relation für Mengen verletzter Constraints



Traditional Constraint Solving

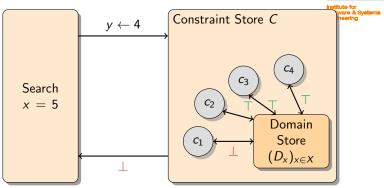




- Eine Kombinationsoperation ∧
- Ein neutrales Element ⊤
- Eine partielle Ordnung $\left(\mathbb{B},\leq_{\mathbb{B}}\right)$ mit $\top<_{\mathbb{R}}\perp$

Klassisches Constraint-Solving

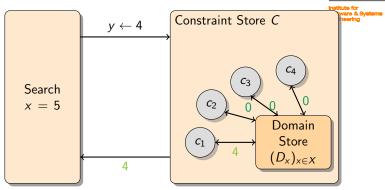




- Eine Kombinationsoperation ∧
- Ein neutrales Element ⊤
- • Eine partielle Ordnung ($\mathbb{B},\leq_{\mathbb{B}}$) mit $\top<_{\mathbb{R}}\perp$

Soft-Constraint-Solving





- Eine Menge von Erfüllungsgraden, e.g., $\{0, \ldots, k\}$
- Eine Kombinationsoperation +
- Ein neutrales Element 0
- • Eine partielle Ordnung (\mathbb{N}, \geq) mit 0 als Top

SoftConstraints in MiniZinc

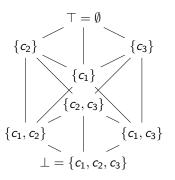


```
% X: \{x,y,z\} D_i = \{1,2,3\}, i in X
% * c1: x + 1 = y * c2: z = y + 2 * c3: x + y <= 3
% (c) ISSE
% isse.uni-augsburg.de/en/software/constraint-relationships/
include "soft_constraints/minizinc_bundle.mzn";
var 1..3: x; var 1..3: y; var 1..3: z;
% read as "soft constraint c1 is satisfied iff x + 1 = y"
constraint x + 1 = y <-> satisfied[1];
constraint z = y + 2 <-> satisfied[2];
constraint x + y <= 3 <-> satisfied[3];
% soft constraint specific for this model
nScs = 3; nCrEdges = 2;
crEdges = [| 2, 1 | 3, 1 |]; % read c2 is less important than c1
solve minimize penSum; % minimize the sum of penalties
```

Search types



The whole valuation space (partially ordered)

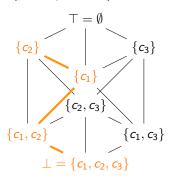


```
%
% Typical Optimization Routine (Branch and Bound):
%
% 1. Look for the first feasible solution
% 2. Impose restrictions on the next feasible solution
```

Search types: Strictly better



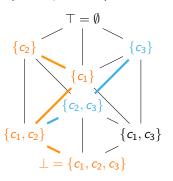
The whole valuation space (partially ordered)



Search types: Only not dominated



The whole valuation space (partially ordered)



Case Studies



Applied to domains where

- Certain properties should really capture preferences, not constraints
- at design time, it is unclear whether an instance is actually solvable
- Solution space is combinatorial
 - Discrete choices
 - Additional hard constraints

Illustrative case studies (found in example-problems)

- Mentor Matching
- Exam Scheduling
- Power Plant Scheduling

Mentor Matching



Goal: Assign mentees (e.g. students) to mentors (e.g. companies) such that

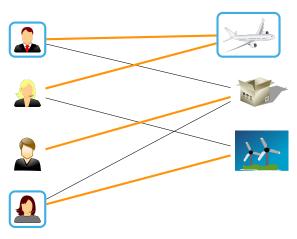
- Students are most satisfied with their mentors
- Companies are satisfied with their mentees
- Two-sided preferences

So far, sounds like a typical stable matching problem, but:

- We do not have a 1:1 mapping (companies advise several students)
- Additional constraints are present
 - Each company has to advise at least I, at most u students
 - The number of advised students should be roughly equal per company (fairness)
 - Students actually despising a company should not be forced to go there (hard exclusion of solutions)

Mentor Matching: Example





This assignment respects the students' preferences (edges) but ignores the companies' preferences. ok, it's not really a matching since companies supervise more than one student...

Mentor Matching: Model



```
int: n; set of int: STUDENT = 1..n;
int: m; set of int: COMPANY = 1..m;
% assign students to companies
array[STUDENT] of var COMPANY: worksAt;
% insert relationships of students and companies here
int: minPerCompany = 2; int: maxPerCompany = 3;
constraint global_cardinality_low_up (
          worksAt, [c | c in COMPANY],
          [minPerCompany | c in COMPANY],
          [maxPerCompany | c in COMPANY]);
solve
:: int_search([ satisfied[mostImpFirst[i]] | i in SOFTCONSTRAINTS],
 input_order, indomain_max, complete)
minimize penSum;
```

Mentor Matching: Preferences



```
n = 3: m = 3:
int: brenner = 1;
int: teufel = 2;
int: fennek = 3;
int: cupgainini = 1;
int: gsm = 2;
int: junedied = 3;
% specify soft constraints, order by relationship
constraint worksAt[teufel] = junedied <-> satisfied[teufJune];
constraint worksAt[teufel] = cupgainini <-> satisfied[teufCap];
constraint worksAt[teufel] = gsm <-> satisfied[teufGsm];
constraint worksAt[fennek] in {cupgainini, gsm} <-> satisfied[fenFavs];
constraint worksAt[fennek ] in {junedied} <-> satisfied[fenOK];
crEdges = [| teufGsm, teufCap | teufGsm, teufJune
           | fenOK, fenFavs |];
```

Mentor Matching: Refinements



Split company and student preferences:

```
% first, our students' preferences
var int: penStud = sum(sc in 1..lastStudentPref)
        (bool2int(not satisfied[sc]) * penalties[sc]);
% now companies' preferences
var int: penComp = sum(sc in lastStudentPref+1..nScs)
        (bool2int(not satisfied[sc]) * penalties[sc]);
```

Optimize lexicographically

```
solve
:: int_search([ satisfied[mostImpFirst[i]] | i in SOFTCONSTRAINTS],%...
%search minimize_lex([penStud, penComp]) /\ if % ...
search minimize_lex([penComp, penStud]) /\ if % ...
```

Mentor Matching: Priority Example



Taken from example: student-company-matching.mzn

```
solve
:: int_search([ satisfied[mostImpFirst[i]] | i in SOFTCONSTRAINTS],%...
search minimize_lex([penStud, penComp]) /\ if %...
```

```
solve
:: int_search([ satisfied[mostImpFirst[i]] | i in SOFTCONSTRAINTS],%...
search minimize_lex([penStud, penComp]) /\ if %...
```

Here, company 1 (cupgainini) wanted to have student 3, and company 2 (APS) did not have any preferences whatsoever (so accepted student 4 instead of 3). Student 4 would have liked company 3 (junedied) better, though.

Mentor Matching: Real Instance



Collected data from winter term

Example

"the favorites":

1. JuneDied-Lynx- HumanIT

2. Cupgainini

"I could live with that":

3. Seamless-German

4. gsm systems

5. Yiehlke

"I think, we won't be happy":

6. APS

7. Delphi Databases

Mentor Matching: Real Instance



- Gave precedence to students
 - After all, what should companies do with unhappy students?
- Search space: 7 companies for 16 students \rightarrow 7¹⁶ = 3.3233 \cdot 10¹³
- Led to a constraint problem with
 - 77 student preferences (soft constraints) from 16 students
 - of a total of 114 soft constraints (37 company preferences)
- Proved optimal solution
 - 4 minutes compilation
 - another 2m 12s solving time

Exam Scheduling



Goal: Assign exam dates to students such that

- Each student likes their appoints (approves of it)
- The number of distinct dates is minimized (to reduce time investment of teachers)



At least 3 options have to be selected

		Approve	Absolutely not
12 February 2016	Morning	0	0
12 February 2016	Afternoon	0	0
18 February 2016	Morning	0	0
18 February 2016	Afternoon	0	0
		0	0
	Name		

- No preference of any student should be weighted higher than another one's
- Solution (exam schedule) is a shared decision

Exam Scheduling: Core Model



See exam-scheduling-approval.mzn:

```
% Exam scheduling example with just a set of
% approved dates and *impossible* ones
include "globals.mzn";
include "soft_constraints/soft_constraints.mzn";
int: n; set of int: STUDENT = 1..n;
int: m; set of int: DATE = 1..m;
array[STUDENT] of set of DATE: possibles;
array[STUDENT] of set of DATE: impossibles;
% the actual decisions
array[STUDENT] of var DATE: scheduled;
int: minPerSlot = 0; int: maxPerSlot = 4;
constraint global_cardinality_low_up(scheduled % minPerSlot, maxPerSlot
constraint forall(s in STUDENT) (not (scheduled[s] in impossibles[s]));
```

Exam Scheduling: Preferences



See exam-scheduling-approval.mzn:

```
% have a soft constraint for every student
nScs = n:
penalties = [ 1 | n in STUDENT]; % equally important in this case
constraint forall(s in STUDENT) (
    (scheduled[s] in possibles[s]) <-> satisfied[s] );
var DATE: scheduledDates;
% constrains that "scheduledDates" different
% values (appointments) appear in "scheduled"
constraint nvalue(scheduledDates, scheduled);
% search variants
solve
:: int_search(satisfied, input_order, indomain_max, complete)
search minimize_lex([scheduledDates, violateds]); % pro teachers
%search minimize_lex([violateds, scheduledDates]); % pro students
```

Exam Scheduling: Real Instance



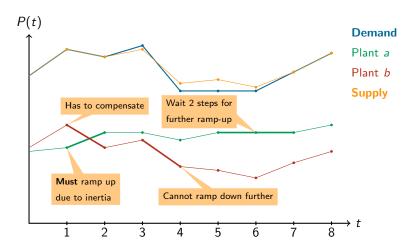
- Collected preferences of 33 students
- over 12 possible dates (6 days, morning and afternoon)
 - Approval set
 - Impossible set
- Aggregated via approval voting (has nice voting-theoretical properties!)
- At most 4 per appointment
- Immediately (61 msec) found an optimal solution that
 - Is approved by every student
 - Is achieved with the minimal number of 9 dates
- Used Strategy:

search minimize_lex([violateds, scheduledDates]); % pro students

Power Plant Scheduling



Goal: Schedule plants such that they meet the demand; See unitCommitment.mzn



Power Plant Scheduling: Core Model



```
include "soft_constraints/soft_constraints_noset.mzn";
include "soft_constraints/cr_types.mzn";
include "soft_constraints/cr_weighting.mzn";
% ground penalties using the appropriate weighting
penalties = [weighting(s, SOFTCONSTRAINTS, crEdges, true)
               | s in SOFTCONSTRAINTS]:
int: T = 5; set of int: WINDOW = 1..T;
array[WINDOW] of float: demand = [10.0, 11.3, 15.2, 20.7, 19.2];
int: P = 3; set of int: PLANTS = 1..P;
array[PLANTS] of float: pMin = [12.0, 5.0, 7.3];
array[PLANTS] of float: pMax = [15.0, 11.3, 9.7];
array[WINDOW, PLANTS] of var 0.0..15.0: supply;
var float: obj;
constraint obj = sum(w in WINDOW) ( abs( sum(p in PLANTS)
          (supply[w, p]) - demand[w]));
```

Power Plant Scheduling: Soft Constraints



```
% ground penalties using the appropriate weighting
penalties = [weighting(s, SOFTCONSTRAINTS, crEdges, true)
               | s in SOFTCONSTRAINTS]:
[...]
% some soft constraints
constraint supply[1, 2] >= 6.0 <-> satisfied[1];
constraint supply[2, 2] >= 6.0 <-> satisfied[2];
% constraint time step 1 seems more urgent
nCrEdges = 1;
crEdges = [| 2, 1 |];
% could do something more sophisticated here
solve minimize obj + penSum;
```

→ Library works with MIP (*Mixed Integer Programming*) as well!

Language Features



а

Language Features: Suitable Weighting



а

Language Features: Consistency Checks



a

Language Features: Variable Ordering



a

Language Features: Redundant Constraints



a

Language Features: Custom Search



а

References I



- Baptiste, P., Laborie, P., Le Pape, C., and Nuijten, W. (2006). Constraint-based scheduling and planning. *Handbook of Constraint Programming*, 22:759–798.
- Barták, R., Salido, M. A., and Rossi, F. (2010). Constraint satisfaction techniques in planning and scheduling. Journal of Intelligent Manufacturing, 21(1):5–15.
- Barták, R. (2003).

Constraint-based scheduling: An introduction for newcomers. In *In Intelligent Manufacturing Systems 2003*, pages 69–74. IFAC Publications, Elsevier Science.

Gadducci, F., Hölzl, M., Monreale, G., and Wirsing, M. (2013).
Soft constraints for lexicographic orders.
In Castro, F., Gelbukh, A., and González, M., editors, *Proc. 12th Mexican Int. Conf. Artificial Intelligence (MICAI'2013)*, Lect. Notes Comp. Sci. 8265, pages 68–79. Springer.

References II



Rossi, F., Beek, P. v., and Walsh, T. (2006).

Handbook of Constraint Programming (Foundations of Artificial Intelligence).

Elsevier Science Inc., New York, NY, USA.

Russell, S. and Norvig, P. (2010).

Artificial Intelligence: A Modern Approach.

Prentice Hall series in artificial intelligence. Prentice Hall.

Schiendorfer, A., Steghöfer, J.-P., Knapp, A., Nafz, F., and Reif, W. (2013). Constraint Relationships for Soft Constraints.

In Bramer, M. and Petridis, M., editors, *Proc. 33rd SGAI Int. Conf. Innovative Techniques and Applications of Artificial Intelligence (Al'13)*, pages 241–255. Springer.

Tsang, E. (1993).

Foundations of constraint satisfaction, volume 289.

Academic press London.